

# class10

Xinlong Wan

## Mini-project

```
candy_file <- "https://raw.githubusercontent.com/fivethirtyeight/data/master/candy-power-ratings.csv"

candy = read.csv(candy_file, row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisp	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0

	hard	bar	pluribus	sugar	percent	price	percent	win	percent
100 Grand	0	1	0		0.732		0.860	66.97	173
3 Musketeers	0	1	0		0.604		0.511	67.60	294
One dime	0	0	0		0.011		0.116	32.26	109
One quarter	0	0	0		0.011		0.511	46.11	650
Air Heads	0	0	0		0.906		0.511	52.34	146
Almond Joy	0	1	0		0.465		0.767	50.34	755

Q1.

```
print(paste(nrow(candy), 'candy types'))
```

[1] "85 candy types"

Q2.

```
print(paste(nrow(candy[candy$fruity==1,]), 'fruity'))
```

[1] "38 fruity"

Q3.

```
print(paste('One dime, ', round(candy["One dime", ]$winpercent,2), '%'))
```

[1] "One dime, 32.26 %"

Q4.

```
print(paste('Kit Kat, ', round(candy["Kit Kat", ]$winpercent,2), '%'))
```

[1] "Kit Kat, 76.77 %"

Q5.

```
print(paste('Tootsie Roll Snack Bars, ', round(candy["Tootsie Roll Snack Bars", ]$winpercent,2), '%'))
```

[1] "Tootsie Roll Snack Bars, 49.65 %"

```
library("skimr")
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

**Variable type: numeric**

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

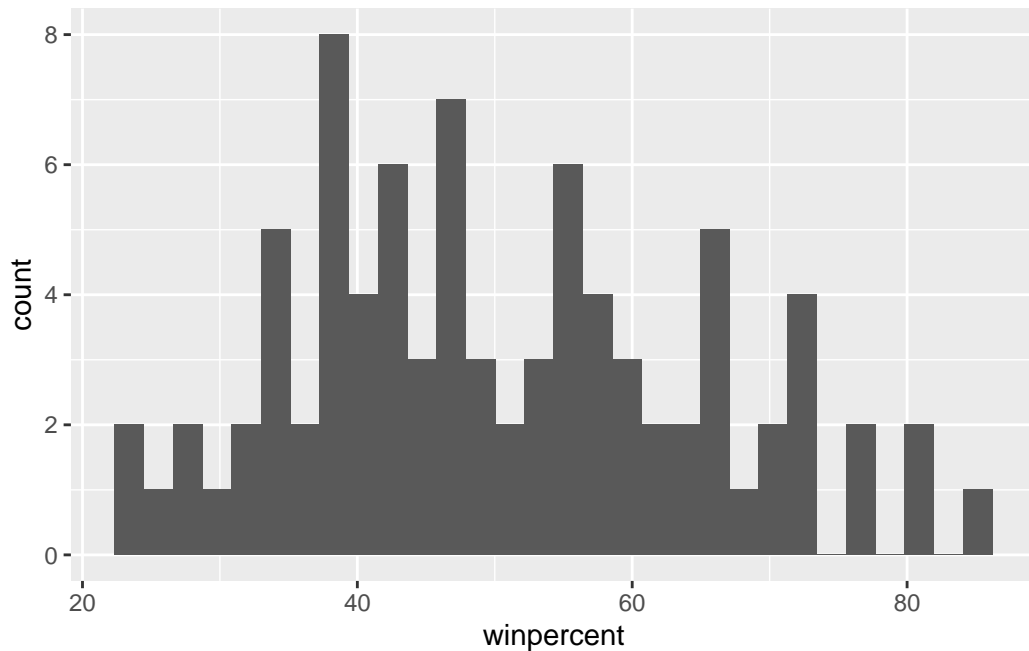
Q6. The sugarpercent, pricepercent, and winpercent columns contains continuous values that are not only 1 and 0. This is because they represent a percentage instead a Yes/No.

Q7. Zero: A specific candy type doesn't contain chocolate. One: A specific candy type contains chocolate.

Q8.

```
library(ggplot2)
ggplot(candy, aes(x=winpercent)) +
  geom_histogram()
```

``stat_bin()` using `bins = 30`. Pick better value with `binwidth`.`



Q9. It's not entirely symmetrical. The distribution skew to the right.

Q10.

```
mean(candy$winpercent) > median(candy$winpercent)
```

```
[1] TRUE
```

The center of the distribution is above 50% as it's skewed to the right.

Q11.

```
mean(candy$winpercent[as.logical(candy$chocolate)]) > mean(candy$winpercent[as.logical(candy$fruit)])
```

```
[1] TRUE
```

On average chocolate candy is higher ranked than fruit candy in terms of winpercentage/

Q12.

```
t.test(candy$winpercent[as.logical(candy$chocolate)], candy$winpercent[as.logical(candy$fruit)])
```

Welch Two Sample t-test

```
data: candy$winpercent[as.logical(candy$chocolate)] and candy$winpercent[as.logical(candy$flavor)]
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

p\_value < 0.001. The difference is statistically significant at alpha=0.001 level.

Q13.

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
candy %>%
  arrange(winpercent) %>%
  head(5) %>%
  rownames()
```

```
[1] "Nik L Nip"           "Boston Baked Beans" "Chiclets"
[4] "Super Bubble"       "Jawbusters"
```

Q14.

```
candy %>%
  arrange(desc(winpercent)) %>%
  head(5) %>%
  rownames()
```


```
[1] "Reese's Peanut Butter cup" "Reese's Miniatures"
[3] "Twix"                      "Kit Kat"
[5] "Snickers"
```

I prefer using arrange, as it allows us to build step by step pipeline that are easier to remember and read.

Q15.

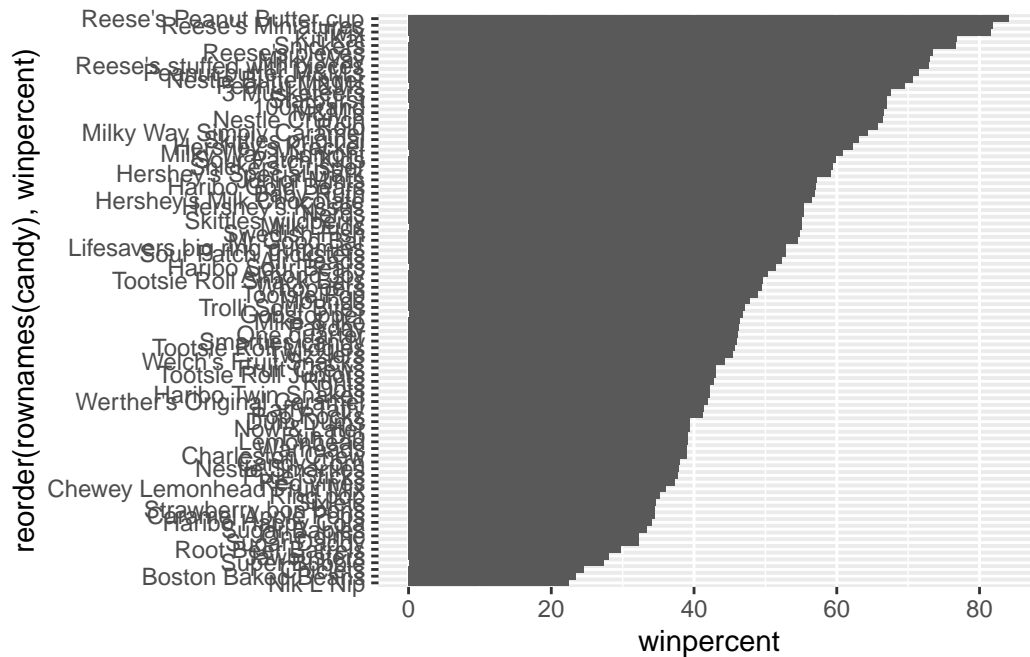
```
library(ggplot2)

ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



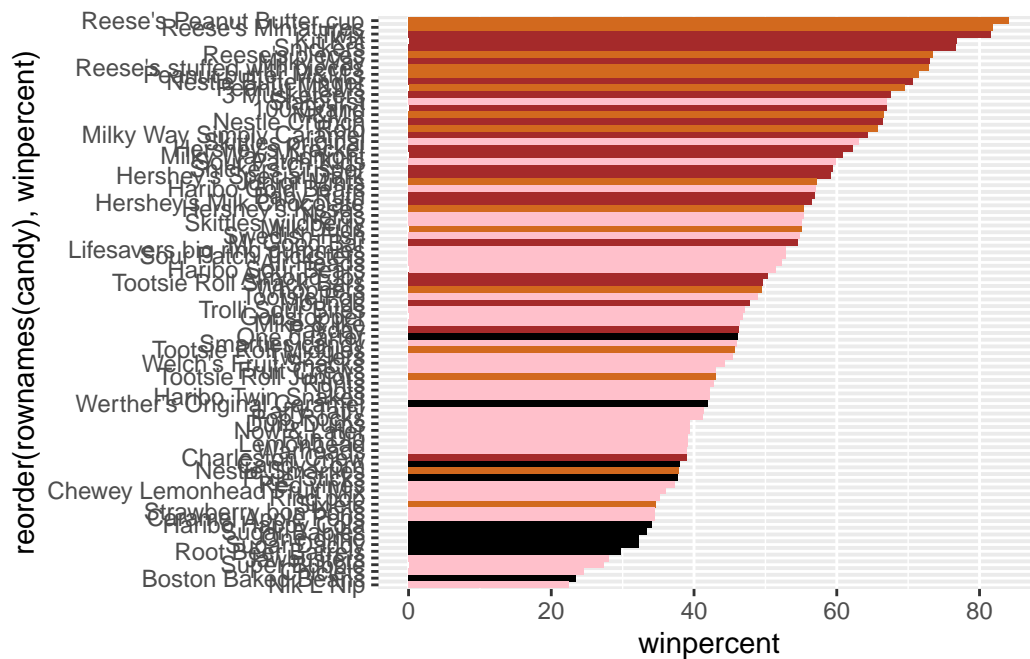
Q16

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```



```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"

ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols)
```



Q17. Sixlets


Q18. Starburst

```
library(ggrepel)

# How about a plot of price vs win
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps





Q19.

```
candy %>%
  arrange(desc(pricepercent)) %>%
  head(5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Nestle Smarties	1	0	0		0	0
Ring pop	0	1	0		0	0
Hershey's Krackel	1	0	0		0	0
Hershey's Milk Chocolate	1	0	0		0	0
	crisped	rice	wafer	hard bar	pluribus	sugarpercent
Nik L Nip		0	0	0	1	0.197
Nestle Smarties		0	0	0	1	0.267
Ring pop		0	1	0	0	0.732
Hershey's Krackel		1	0	1	0	0.430
Hershey's Milk Chocolate		0	0	1	0	0.430
	pricepercent	winpercent				
Nik L Nip	0.976	22.44534				
Nestle Smarties	0.976	37.88719				
Ring pop	0.965	35.29076				

Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Reese's Miniatures cost the least among the top five popular candies.

Q20.

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

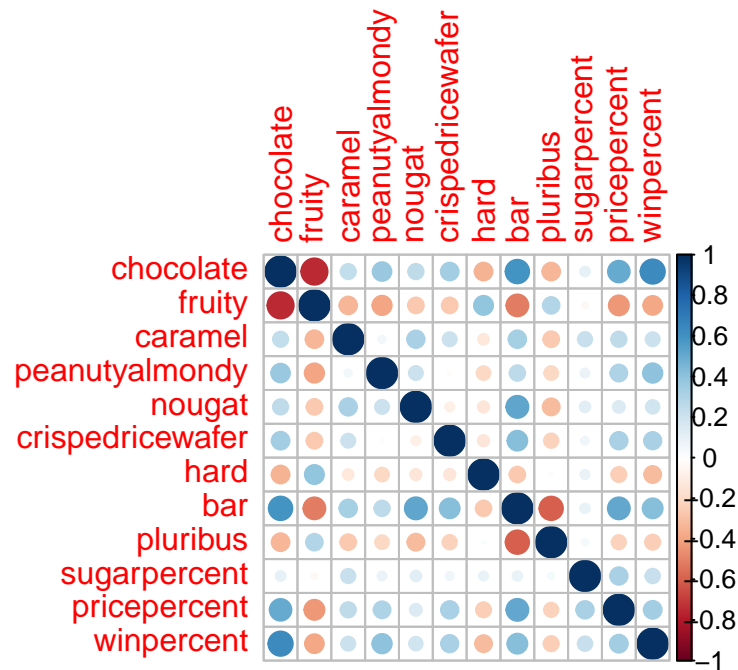
	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Nik L Nip is the least popular.

```
library(corrplot)
```

corrplot 0.92 loaded

```
cij <- cor(candy)
corrplot(cij)
```



Q22. fruity and chocolate are anti-correlated. Bar and pluribus are very highly negatively negatively correlated. Many others are slightly negatively correlated.

Q23. winpercentage and chocolate. Bar and chocolate.

```
pca <- prcomp(candy, scale=FALSE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	14.7231	0.70241	0.47762	0.37292	0.34641	0.33614	0.30748
Proportion of Variance	0.9935	0.00226	0.00105	0.00064	0.00055	0.00052	0.00043
Cumulative Proportion	0.9935	0.99574	0.99678	0.99742	0.99797	0.99849	0.99892

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.27417	0.23826	0.21435	0.18434	0.15331
Proportion of Variance	0.00034	0.00026	0.00021	0.00016	0.00011
Cumulative Proportion	0.99927	0.99953	0.99974	0.99989	1.00000

```
pca$rotation[,1]
```

chocolate                      fruity                      caramel      peanutyalmondy

-0.021594409	0.012979597	-0.005422631	-0.010306565
nougat	crispedricewafer	hard	bar
-0.003755535	-0.006106095	0.008097109	-0.012699628
pluribus	sugarpercent	pricepercent	winpercent
0.008474153	-0.004401607	-0.006717124	-0.999407346

PC1 is dominated by winpercentage, since it's on a very different scale.

We should scale it.

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

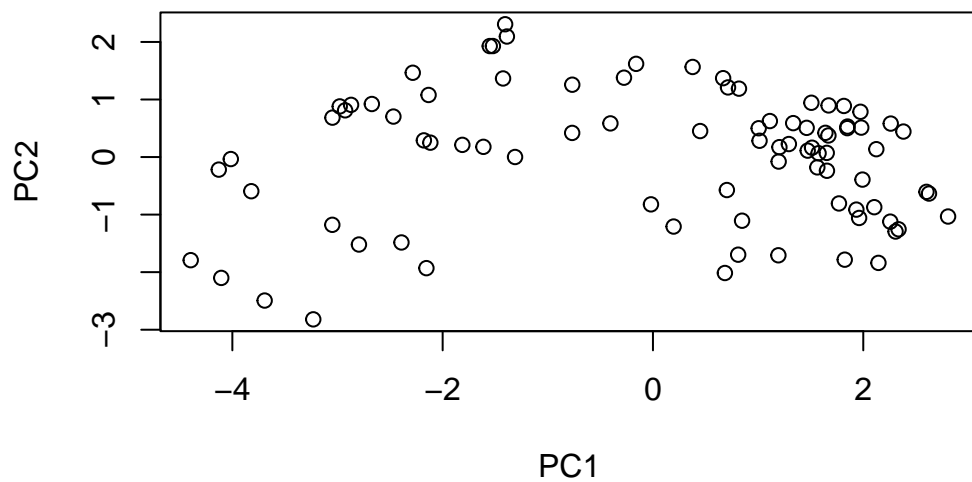
Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369


  

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

```
plot(pca$x[,1:2])
```




```
plot(pca$x[,1:2], col=my_cols, pch=16)
```



```
# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])
```

```
p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)
```

p




```
p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown)",
       caption="Data from 538")
```

Warning: ggrepel: 59 unlabeled data points (too many overlaps). Consider increasing max.overlaps

## Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown),



Data from 538

If you want to see more candy labels you can change the `max.overlaps` value to allow more overlapping labels or pass the ggplot object `p` to `plotly` like so to generate an interactive plot that you can mouse over to see labels:

```
library(plotly)
```

Attaching package: 'plotly'

The following object is masked from 'package:ggplot2':

`last_plot`

The following object is masked from 'package:stats':

`filter`


The following object is masked from 'package:graphics':

`layout`

```
# ggplotly(p)
```

```
par(mar=c(8,4,2,2))
```

```
barplot(pca$rotation[,1], las=2, ylab="PC1 Contribution")
```



Q24. fruity, hard, pluribus. PC1 has a large positive association with fruity, hard, and pluribus. This makes sense since these three factors are highly correlated. Fruity candy is usually hard and packed together in a box of multiple candies.